

# Database Tuning

## Queries

### 1. Query 1:


1. Description: Find posts of users in city 1 with more than 20 likes.
2. Query:

```
SELECT p.body FROM posts p
JOIN user_posts up ON up.post_id = p.id
JOIN users u ON u.id = up.user_id
JOIN likes l ON l.post_id = p.id
WHERE u.city = 'city 1'
GROUP BY p.id
HAVING COUNT(l.id) > 20;
```

### 3. Explain:

1. The query will first join the **posts** table with **user\_posts** table on **post\_id** column.
2. Then it will join the result with **users** table on **user\_id** column.
3. Then it will join the result with **likes** table on **post\_id** column.
4. Then it will filter the result by **city** column.
5. Then it will group the result by **id** column.
6. Then it will filter the result by **COUNT(l.id) > 20**.

### 4. Query Tree:



### 2. Query 2:

1. Description: Find all the comments and posts of a user that is older than 25 and lives in city 1
2. Query:

```
SELECT uc.comment_id, up.post_id FROM users u
INNER JOIN user_posts up ON u.id = up.user_id
INNER JOIN user_comments uc ON u.id = uc.user_id
WHERE u.age > 25 and u.city = 'city 1'
GROUP BY up.post_id , uc.comment_id;
```


### 3. Explain:

1. The query will first join the **users** table with **user\_posts** table on **id** column.
2. Then it will join the result with **user\_comments** table on **user\_id** column.

3. Then it will filter the result by `age > 25` and `city = 'city 1'`.

4. Then it will group the result by `post_id` and `comment_id`.

4. Query Tree:



3. Query 3:

1. Description: Find the posts with most likes in city 1.


2. Query:

```
SELECT p.body, u.city, COUNT(l.id) AS likes_count
FROM posts p
JOIN user_posts up ON up.post_id = p.id
JOIN users u ON u.id = up.user_id
JOIN likes l ON l.post_id = p.id
WHERE u.city = 'city 1'
GROUP BY u.city, p.body
ORDER BY likes_count DESC;
```

3. Explain:

1. The query will first join the `posts` table with `user_posts` table on `post_id` column.
2. Then it will join the result with `users` table on `user_id` column.
3. Then it will join the result with `likes` table on `post_id` column.
4. Then it will filter the result by `city = 'city 1'`.
5. Then it will group the result by `city` and `body`.
6. Then it will order the result by `likes_count` in descending order.

4. Query Tree:



4. Query 4:

1. Description: Get the full details of all the posts that has 1 in the title and 2 in the body from users that are older than 28 and live in city 1

2. Query:

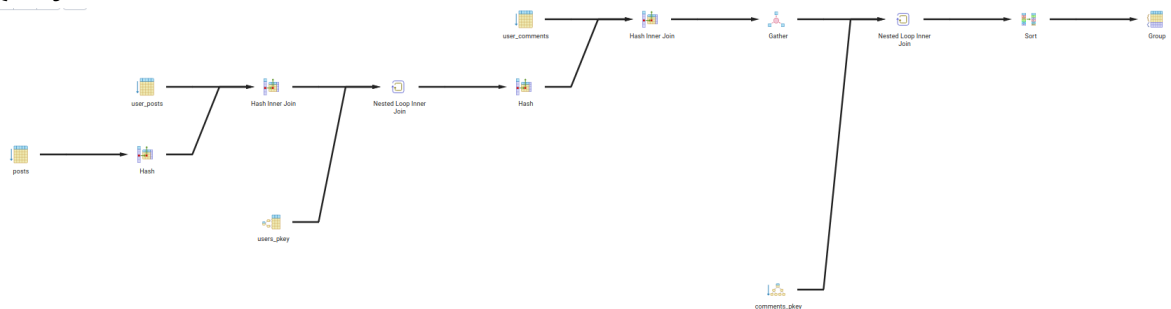
```
SELECT up.post_id , p.body , p.title
FROM users u
INNER JOIN user_posts up ON u.id = up.user_id
INNER JOIN user_comments uc ON u.id = uc.user_id
INNER JOIN posts p ON p.id = up.post_id
INNER JOIN comments c ON c.id = uc.comment_id
WHERE u.age > 28 and u.City = 'city 1' and
```

```
p.title LIKE '%title 1%' and p.body LIKE '%body 2%'  
GROUP BY up.post_id, p.body, p.title;
```

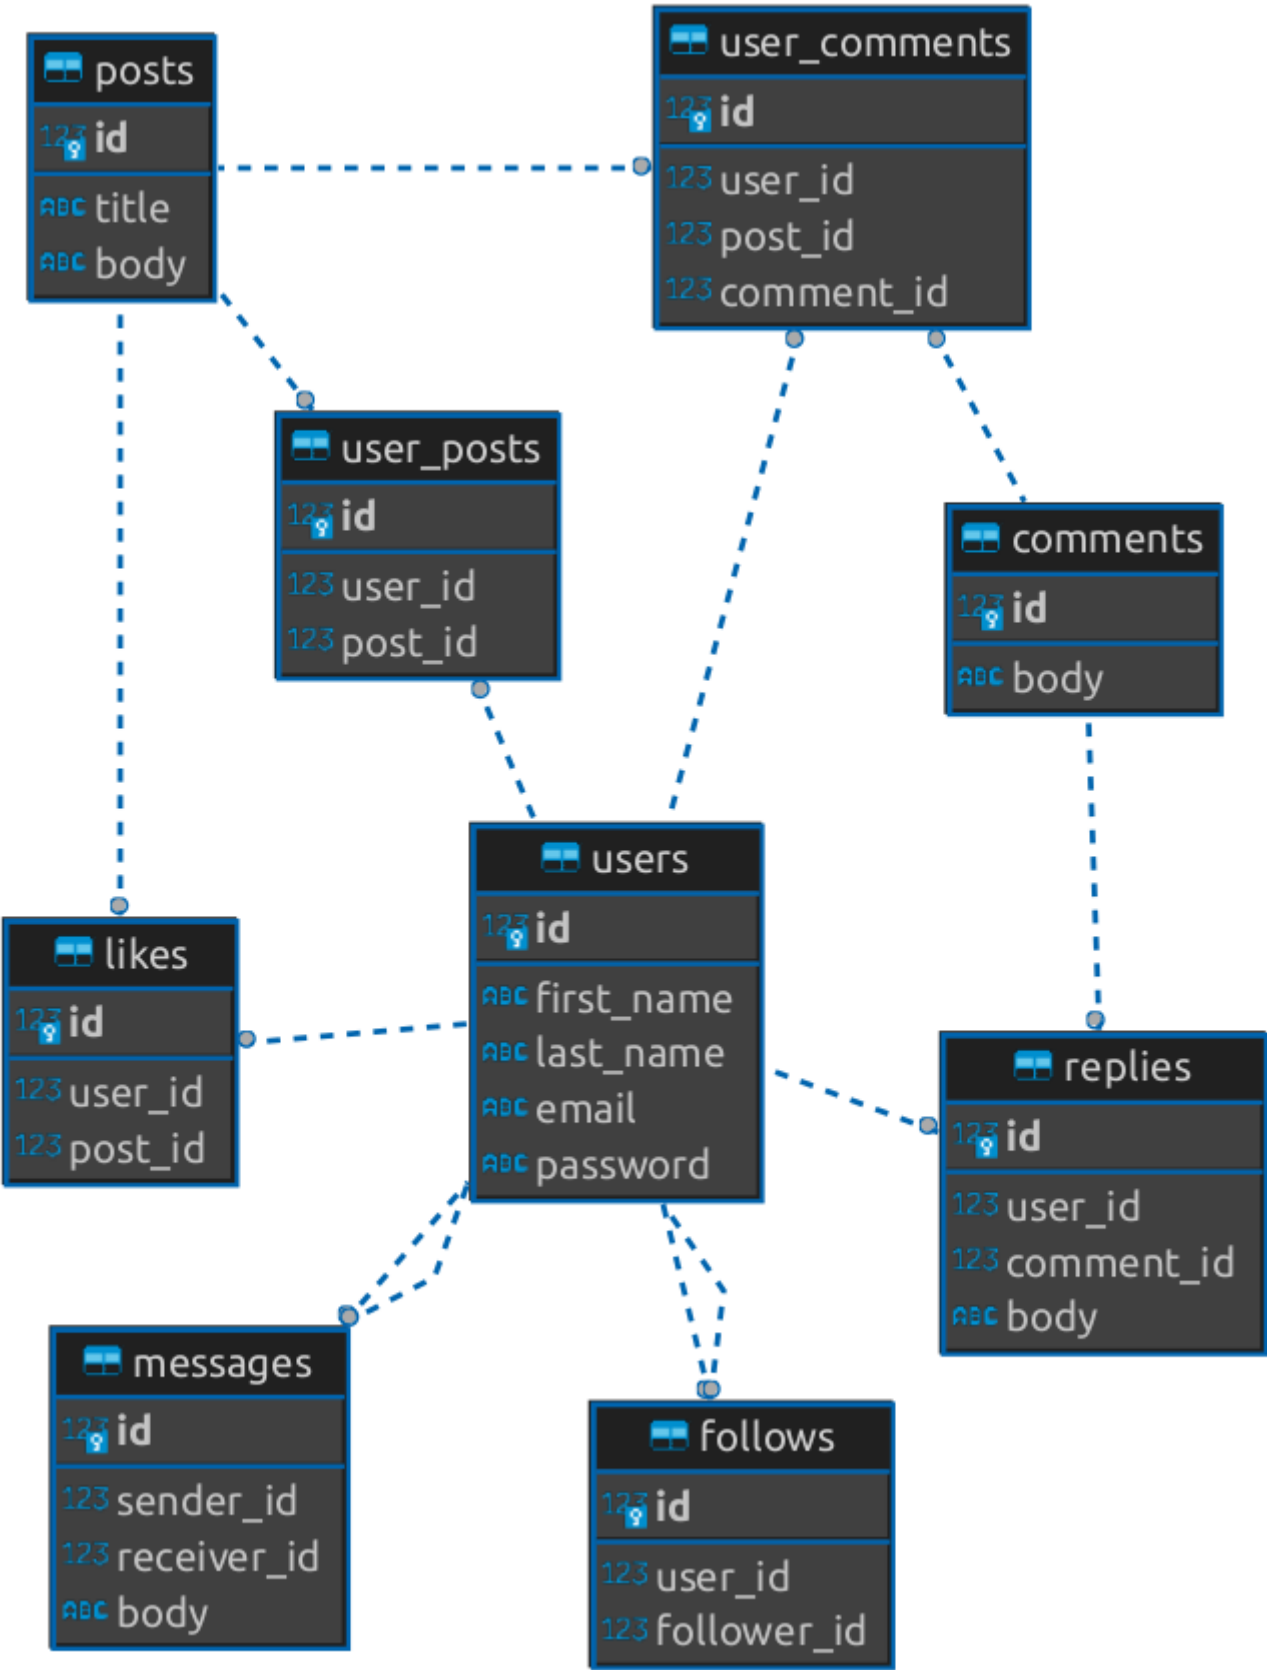
### 3. Explain:

1. The query will first join the **users** table with **user\_posts** table on **id** column.
2. Then it will join the result with **user\_comments** table on **user\_id** column.
3. Then it will join the result with **posts** table on **post\_id** column.
4. Then it will join the result with **comments** table on **comment\_id** column.
5. Then it will filter the result by **age > 28** and **city = 'city 1'** and **title LIKE '%title 1%'** and **body LIKE '%body 2%'**.
6. Then it will group the result by **post\_id, body, title**.

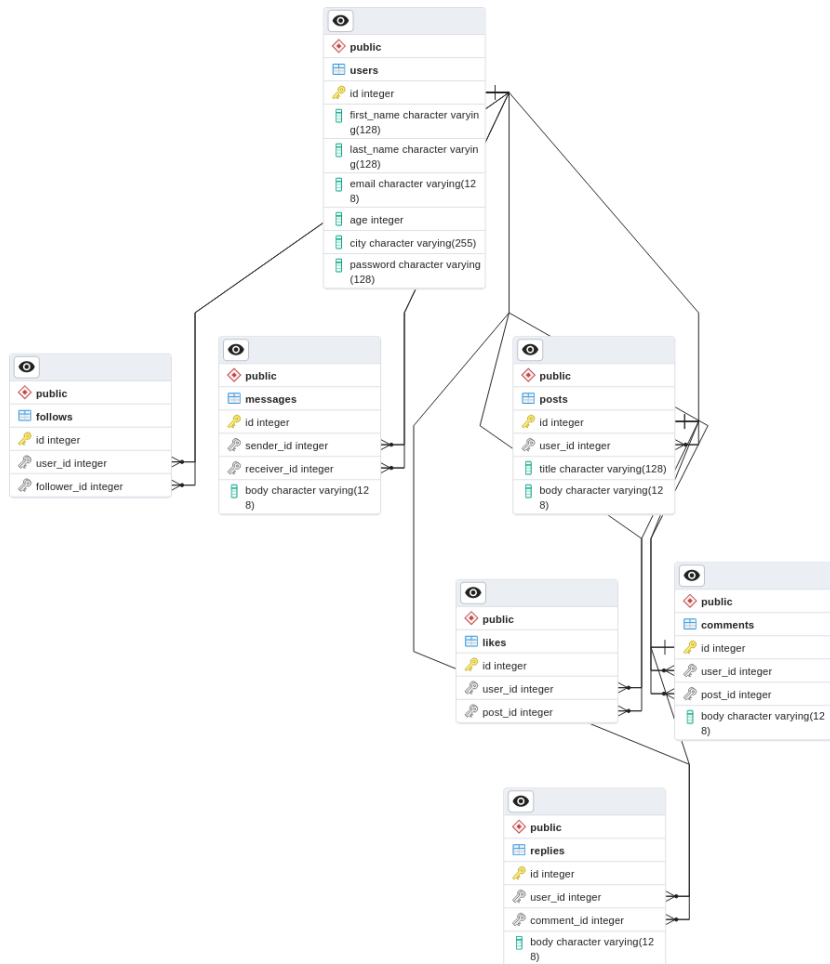
### 4. Query Tree:



## Postgres Server Old Schema



Postgres Server Enhanced Schema



## Optimizations

### Subquery Optimizations

#### 1. Query 1:

1. Description: Find posts of users in city 1 with more than 20 likes.

2. Query:

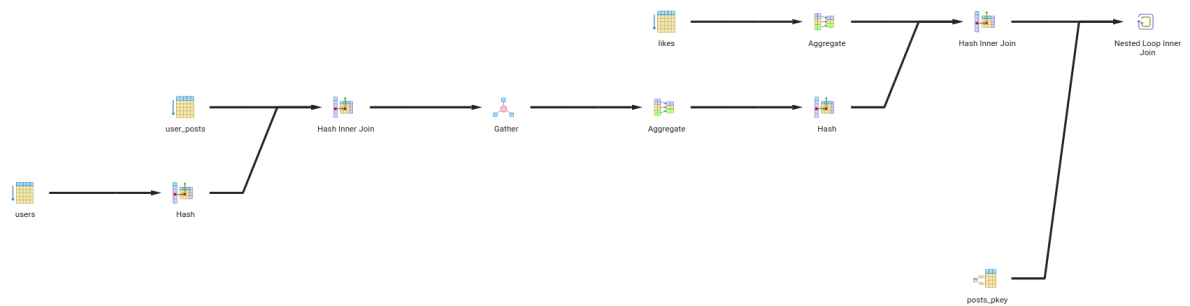
```

SELECT p.body FROM posts p
where p.id in
  (SELECT up.post_id FROM user_posts up JOIN users u ON u.id =
up.user_id WHERE u.city = 'city 1')
And p.id in
  (SELECT l.post_id FROM likes l GROUP BY l.post_id HAVING
COUNT(l.id) > 20);
  
```

#### 3. Explain:

1. The query will first join the **user\_posts** table with **users** table on **id** column.
2. Then it will filter the result by **city = 'city 1'**.
3. Then it will join the result with **likes** table on **post\_id** column.
4. Then it will group the result by **post\_id**.
5. Then it will filter the result by **COUNT(l.id) > 20**.
6. Then it will select the **body** column from the result.

## 4. Query Tree:



## 2. Query 2:

1. Description: Find all the comments and posts of a user that is older than 25 and lives in city 1


2. Query:

```
with CTE as (SELECT * FROM users WHERE age > 25 and city =
'city 1')
SELECT uc.comment_id, up.post_id FROM CTE AS u
INNER JOIN user_posts up ON u.id = up.user_id
INNER JOIN user_comments uc ON u.id = uc.user_id
GROUP BY up.post_id , uc.comment_id;
```

## 3. Explain:

1. The query will first filter the **users** table by **age > 25** and **city = 'city 1'** using a CTE.
2. Then it will join the result with **user\_posts** table on **user\_id** column.
3. Then it will join the result with **user\_comments** table on **user\_id** column.
4. Then it will group the result by **post\_id** and **comment\_id**.

## 4. Query Tree:



## 3. Query 3:

1. Description: Find the posts with most likes in city 1.


2. Query:

```
SELECT p.body, u.city, COUNT(l.id) AS likes_count
FROM posts p
INNER JOIN user_posts up ON up.post_id = p.id
INNER JOIN users u ON u.id = up.user_id
INNER JOIN likes l ON l.post_id = p.id
WHERE u.city = 'city 1'
GROUP BY u.city, p.body
ORDER BY likes_count DESC;
```

## 3. Explain:

1. The query will first join the **posts** table with **user\_posts** table on **post\_id** column.
2. Then it will join the result with **users** table on **user\_id** column.
3. Then it will join the result with **likes** table on **post\_id** column.
4. Then it will filter the result by **city = 'city 1'**.
5. Then it will group the result by **city, body**.
6. Then it will order the result by **likes\_count** in descending order.
7. Then it will select the **body, city, likes\_count** columns.

## 4. Query Tree:



## 4. Query 4:

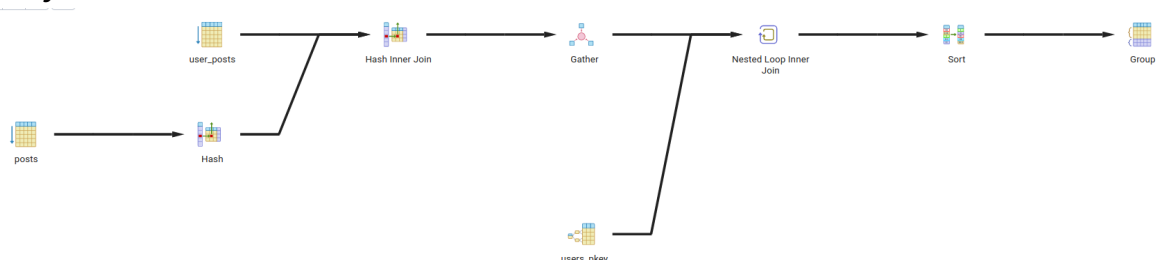
1. Description: Get the full details of all the posts that has 1 in the title and 2 in the body from users that are older than 28 and live in city 1
2. Query:

```
with CTE as (SELECT * FROM users WHERE age > 25 and city = 'city 1')
SELECT up.post_id FROM CTE AS u
INNER JOIN user_posts up ON u.id = up.user_id
inner join posts p on p.id = up.post_id
where p.title LIKE '%title 1%' and p.body LIKE '%body 2%'
GROUP BY up.post_id;
```

## 3. Explain:

1. The query will first filter the **users** table by **age > 25** and **city = 'city 1'** using a CTE.
2. Then it will join the result with **user\_posts** table on **user\_id** column.
3. Then it will join the result with **posts** table on **post\_id** column.
4. Then it will filter the result by **title LIKE '%title 1%'** and **body LIKE '%body 2%'**.
5. Then it will group the result by **post\_id**.
6. Then it will select the **post\_id** column.

## 4. Query Tree:



## Indexes

### 1. Query 1:

1. Description: Find posts of users in city 1 with more than 20 likes.

2. Query:


```
-- add index for post_id in user_posts table if not exists
CREATE INDEX IF NOT EXISTS user_posts_post_id_idx ON
user_posts (post_id);
-- add index for post_id in likes table if not exists
CREATE INDEX IF NOT EXISTS likes_post_id_idx ON likes
(post_id);
-- add index for city in users table if not exists
CREATE INDEX IF NOT EXISTS users_city_idx ON users (city);
-- add index for user_id in user_posts table if not exists
CREATE INDEX IF NOT EXISTS user_posts_user_id_idx ON
user_posts (user_id);

-- use explain analyze to see the difference in execution time
SELECT p.body FROM posts p
JOIN user_posts up ON up.post_id = p.id
JOIN users u ON u.id = up.user_id
JOIN likes l ON l.post_id = p.id
WHERE u.city = 'city 1'
GROUP BY p.id
HAVING COUNT(l.id) > 20;
```

### 3. Explain:

1. We have added indexes for **post\_id** in **user\_posts** table, **post\_id** in **likes** table, **city** in **users** table, and **user\_id** in **user\_posts** table.
2. We have added those indexes to make the query more computationally efficient.
3. For instance, we have added an index for **city** in **users** table to make the filtering by **city = 'city 1'** more efficient.
4. We have added an index for **post\_id** in **user\_posts** table to make the join between **posts** table and **user\_posts** table more efficient.
5. We have added an index for **post\_id** in **likes** table to make the join between **posts** table and **likes** table more efficient.
6. We have added an index for **user\_id** in **user\_posts** table to make the join between **users** table and **user\_posts** table more efficient.

### 4. Query Tree:



### 2. Query 2:

1. Description: Find all the comments and posts of a user that is older than 25 and lives in city 1



## 2. Query:

```
-- add composite index for users table if not exists
CREATE INDEX IF NOT EXISTS users_age_city_idx ON users (city,
age);


-- add index for user_id in user_posts table if not exists
CREATE INDEX IF NOT EXISTS user_posts_user_id_idx ON
user_posts (user_id);

-- we changed the order of the conditions in the where clause
to make indexing more efficient
SELECT uc.comment_id, up.post_id FROM users u
INNER JOIN user_posts up ON u.id = up.user_id
INNER JOIN user_comments uc ON u.id = uc.user_id
WHERE u.city = 'city 1' and u.age > 25
GROUP BY up.post_id , uc.comment_id;
```

## 3. Explain:

1. We have added a composite index for **city** and **age** in **users** table, and an index for **user\_id** in **user\_posts** table.
2. We have added those indexes to make the query more computationally efficient.
3. We have added a composite index for **city** and **age** in **users** table to make the filtering by **city = 'city 1'** and **age > 25** more efficient.
4. We changed the order of the conditions in the **where** clause to make indexing more efficient.
5. We have added an index for **user\_id** in **user\_posts** table to make the join between **users** table and **user\_posts** table more efficient.

## 4. Query Tree:



## 3. Query 3:

1. Description: Find the posts with most likes in city 1.
2. Query:

```
-- add index for city in users table if not exists
CREATE INDEX IF NOT EXISTS users_city_idx ON users (city);
-- add index for user_id in user_posts table if not exists
CREATE INDEX IF NOT EXISTS user_posts_user_id_idx ON
user_posts (user_id);


-- get the posts with max likes in every city
SELECT p.body, u.city, COUNT(l.id) AS likes_count
FROM posts p
JOIN user_posts up ON up.post_id = p.id
JOIN users u ON u.id = up.user_id
```

```
JOIN likes l ON l.post_id = p.id
WHERE u.city = 'city 1'
GROUP BY u.city, p.body
ORDER BY likes_count DESC;
```

### 3. Explain:

1. We have added an index for **city** in **users** table, and an index for **user\_id** in **user\_posts** table.
2. We have added those indexes to make the query more computationally efficient.
3. We have added an index for **city** in **users** table to make the filtering by **city = 'city 1'** more efficient.
4. We have added an index for **user\_id** in **user\_posts** table to make the join between **users** table and **user\_posts** table more efficient.

### 4. Query Tree:



### 4. Query 4:

1. Description: Get the full details of all the posts that has 1 in the title and 2 in the body from users that are older than 28 and live in city 1
2. Query:

```
-- add composite index for users table if not exists
CREATE INDEX IF NOT EXISTS users_age_city_idx ON users (city,
age);

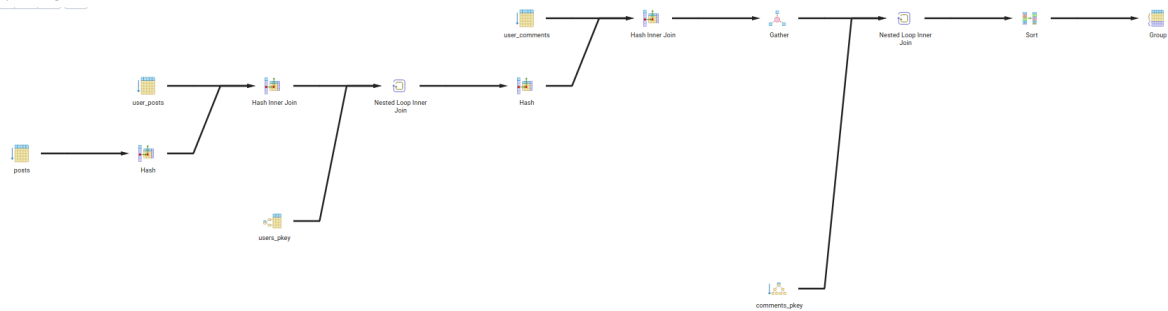
-- we changed the order of the conditions in the where clause
to make indexing more efficient
-- indexing won't be efficient due to the use of LIKE
operator
SELECT up.post_id , p.body , p.title
FROM users u
INNER JOIN user_posts up ON u.id = up.user_id
INNER JOIN user_comments uc ON u.id = uc.user_id
INNER JOIN posts p ON p.id = up.post_id
INNER JOIN comments c ON c.id = uc.comment_id
WHERE u.City = 'city 1' and u.age > 28 and
p.title LIKE '%title 12%' and p.body LIKE '%body 22%'
GROUP BY up.post_id, p.body, p.title;
```

### 3. Explain:

1. We have added a composite index for **city** and **age** in **users** table.
2. We have added those indexes to make the query more computationally efficient.
3. We have added a composite index for **city** and **age** in **users** table to make the filtering by **city = 'city 1'** and **age > 28** more efficient.

4. We changed the order of the conditions in the **where** clause to make indexing more efficient.

4. Query Tree:



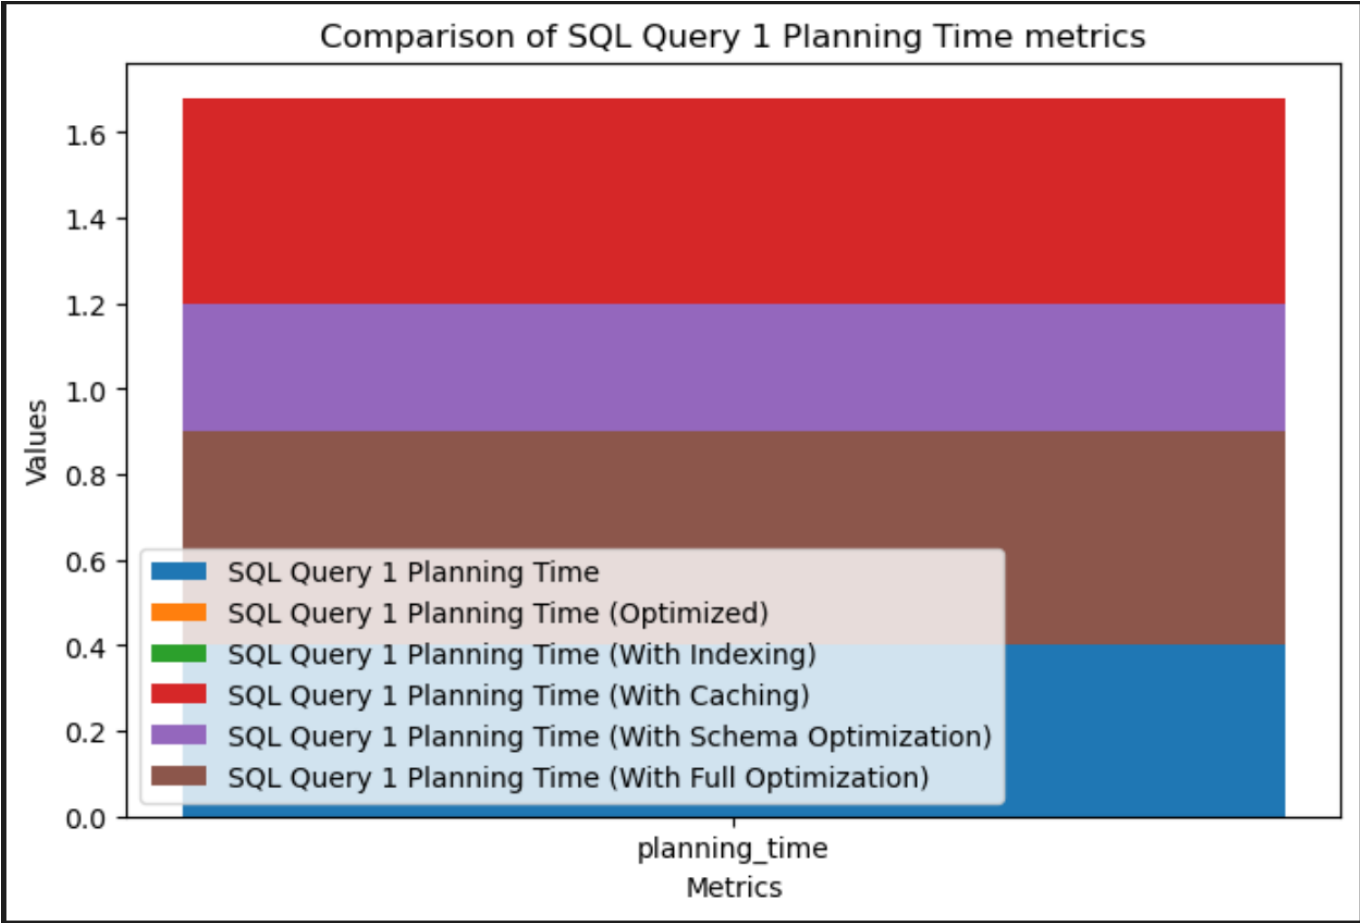

Postgres Server Statistics Report

Table Name	Row Count	Main Key	Indexes	FK	Identity Column	Max Row Size(Bytes)
users	50k	Yes	No	No	Yes	87
posts	50k	Yes	No	No	Yes	50
user_posts	48734	Yes	No	Yes	Yes	35
comments	50k	Yes	No	No	Yes	38
user_comments	48734	Yes	No	Yes	Yes	39
replies	9920	Yes	No	Yes	Yes	46
likes	48734	Yes	No	Yes	Yes	35
follows	39302	Yes	No	Yes	Yes	35
messages	50k	Yes	No	Yes	Yes	46

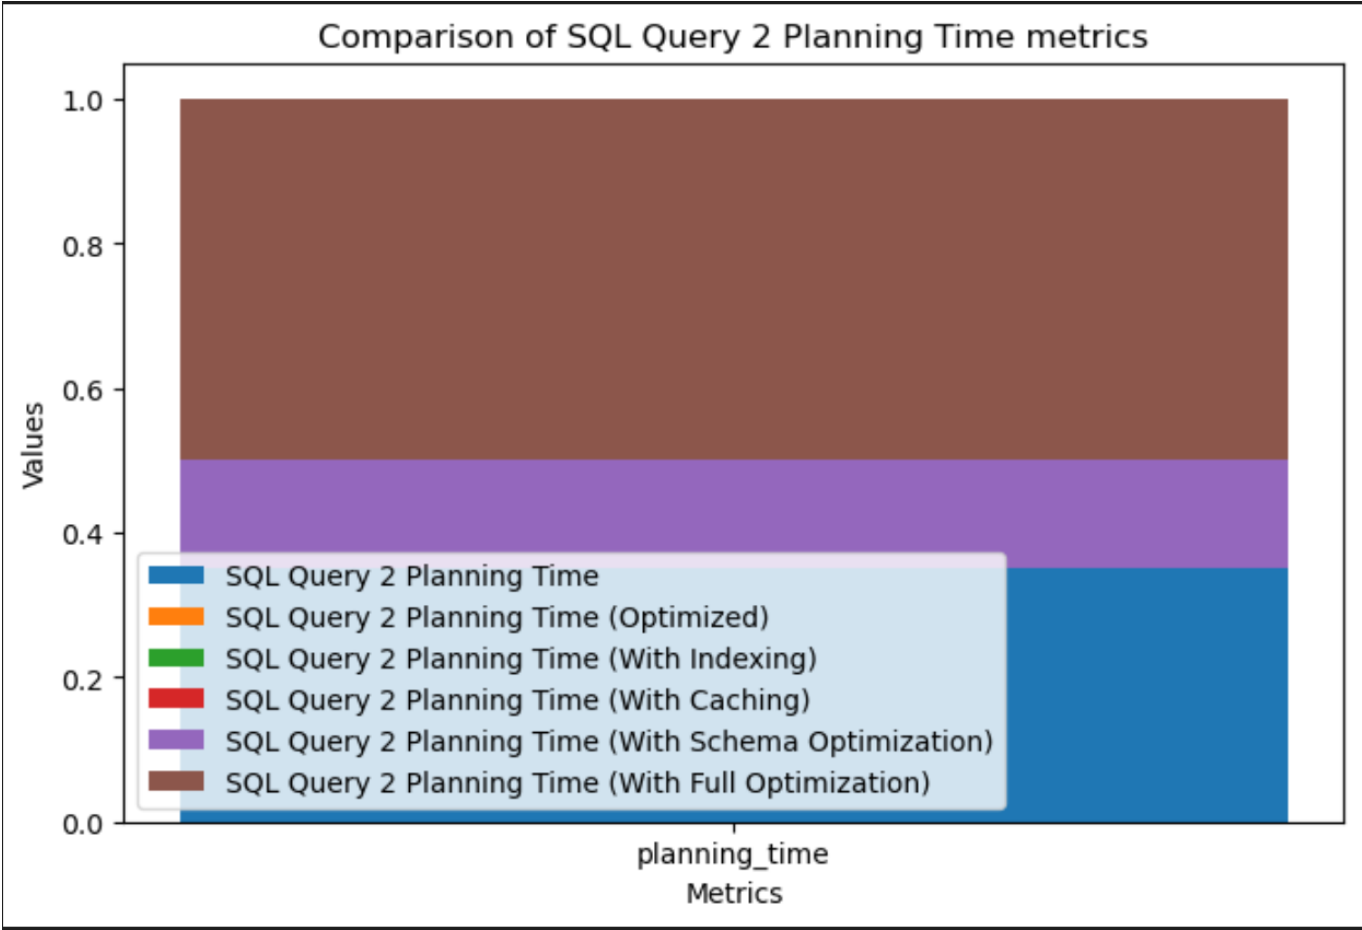

Validation Details and Comparisons

Comparison of the plan and execution times for the queries with different optimization techniques.



Query 1





Query 2



Query 3





Query 4





Comparison of the plan and execution times for the queries with different sizes.

Query 1

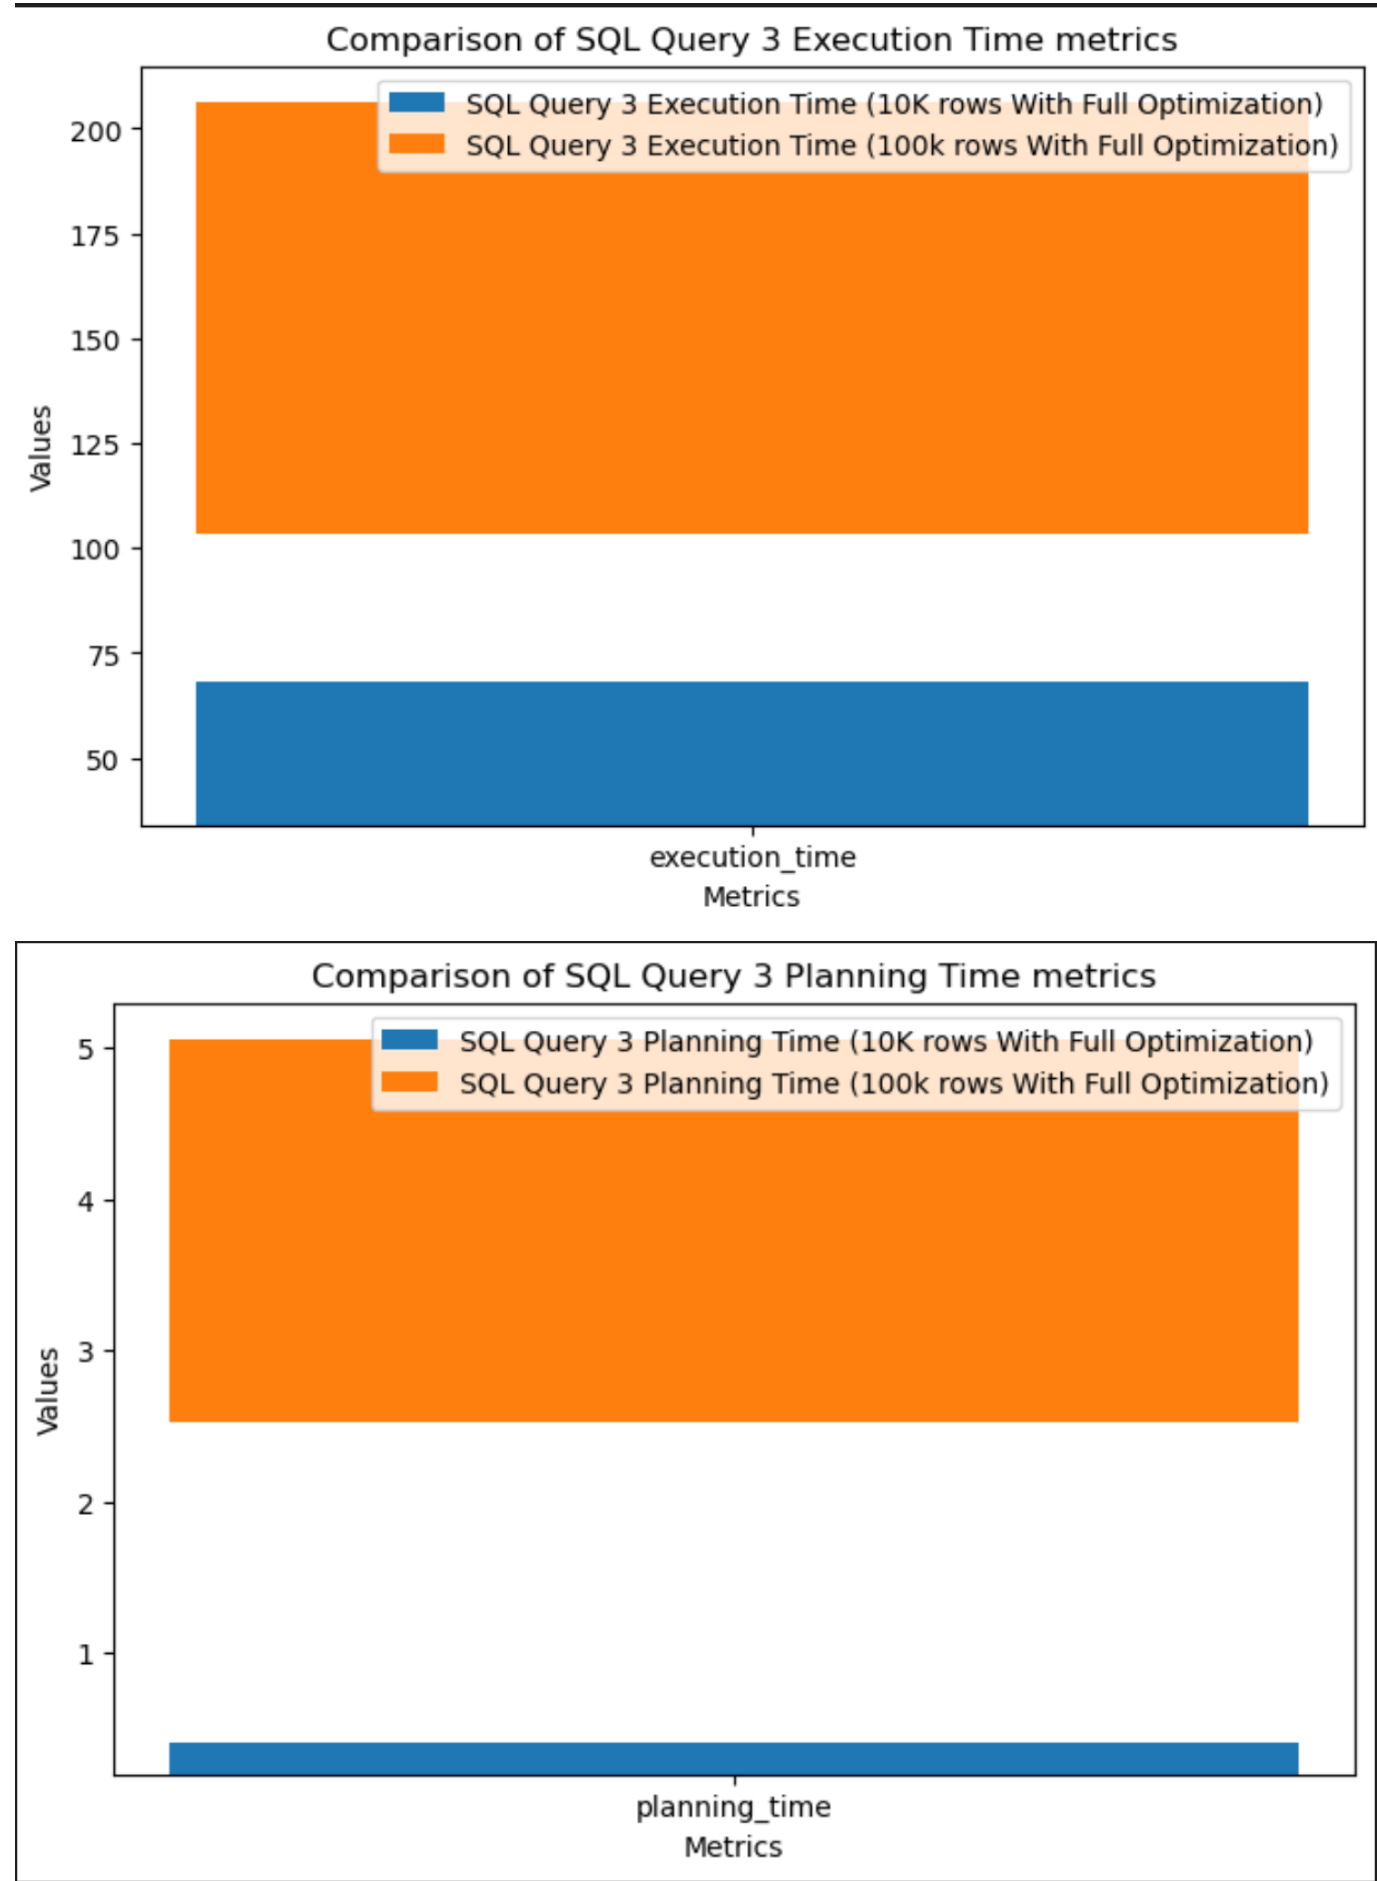


Query 2






Query 3





Query 4





Comparison of the plan and execution times for the queries with different disks(HDD VS. SSD).



Query 1





Query 2



Query 3



Query 4



## System Specifications

Name	Version
------	---------

Name	Version
OS	Ubuntu 22.04 LTS
CPU1	Intel(R) Core(TM) i7-1087H CPU @ 2.20GHz
CPU2	AMD Ryzen™ 7 5800 @ 2.20GHz
Disk1	HDD 1 TB
Disk2	SSD 1 TB
RAM	16 GB

## Developer Guide

### How to run?

1. Run `docker-compose up`.
2. You will see a message whether the creation was `Okay` or `Not Okay`.
3. Exec into the container.
4. Populate the data by running with the required size `./mnt/scripts/populate_data.sh 10k`.
5. Run the query using `./mnt/scripts/run_queries.sh query1`.
6. Don't forget to stop the container after you are done.
7. Run `docker-compose down --volumes` to delete the volumes.

how to copy and paste file in bash script?

```
cp /mnt/scripts/queries/query1.js /mnt/scripts/queries/query1.js.bak
```

### For NOSQL

1. Run `docker-compose up mongodb`
2. Container will create collections upon startup also it will insert dummy data into them
3. Use `connect_db.sh` to exec into container
4. Run the following commands

```
mongosh
use admin
db.auth("root", "pass12345")
use social
```

5. The container starts initially with 10k data for each collection , to change this
  - Open a shell in the container using `connect_db.sh`
  - change the number fo dummy data using `export DUMMY_SIZE=1000000`
  - run `mongosh`
  - Authenticate yourself using step 4 commands



- load the script again using `load("./insert_dummy.js")`

## How can you contribute?

1. Create queries in the `queries` folder.
2. Add the queries to the `run_queries.sh` file.

## How to add data?

1. Add a new script in the `dummyData` folder.
2. Add the new size of the data in the `populate_data.sh` file.

You can find the schema of the database in the `DDL` folder.

## Helpful mongo commands

1. Note that mongo command should be installed on the computer. On Linux this should be install `mongodb-org-shell` package.
2. Connect to MongoDB server `mongo admin -u root -p pass12345` It will connect to localhost port 27017.
3. Show databases: `show dbs`
4. Create new non-existent database: `use mydatabase`
5. Show collections: `show collections`
6. Show contents/documents of a collection: `db.your_collection_name.find()`
7. Save a data to a collection: `db.your_collection_name.save({"name": "Sony AK"})`
8. Show database version: `db.version()`
9. Show database status: `db.stats()`